

GROUP2505 – Log-likelihood analysis for Restricted Boltzmann Machines

Elias Maria Bonasera, Alberto Casellato, Nicola Garbin, and Francesco Pazzocco
(Dated: March 27, 2025)

[illegible]

1. INTRODUCTION

The main goal of this assignment is to explore how the performance of an RBM changes for different choices of the hyperparameters of the model, using the MNIST digits as the database; in particular, using the log-likelihood evaluation, we explore the trend of the model during the learning process as the number of epochs increases.

2. RESTRICTED BOLTZMANN MACHINES

Restricted Boltzmann Machines (RBM) are a powerful kind of generative models designed to accomplish training processes relatively fast. In RBMs, a set of D binary visible units i of state $\{v_i\}_{i=1,\dots,D}$ is symmetrically connected to a set of L binary hidden units μ of state $\{h_\mu\}_{\mu=1,\dots,L}$; the continuous weight $w_{i\mu}$ quantifies the strength between unit i and unit μ (see Figure 1). RBMs use an energy function to define the probability distribution over the input data [1] [2].

In the training process the energy of the configuration is minimized by adjusting the parameters θ . One of the most adopted training algorithm is Contrastive Divergence which allows to approximate the gradient of the likelihood to update the parameters. During this process, a cyclic Gibbs sampling is performed setting the visible units given the hidden ones and vice versa, according to the following probabilities p :

$$p(h_\mu = 1 \mid \mathbf{v}) = \sigma(b_\mu + \sum_i v_i w_{i\mu}) \quad (1)$$

$$p(v_i = 1 \mid \mathbf{h}) = \sigma(a_i + \sum_{\mu} h_{\mu} w_{i\mu}) \quad (2)$$

where $\sigma(x) = 1/(1 + e^{-x})$ is the logistic sigmoid function, a_i is the bias of the i -th visible unit and b_μ is the bias of the μ -th hidden unit; they act shifting the sigmoid function $\sigma(x)$. The absence of links among units of the same type simplifies the training process. Moreover, the number of iterations of Eq. 1 and Eq. 2 can be set to 2 if real data is used to fix $\{v_i\}_{i=1,\dots,D}$ in the first place.




FIG. 1. A Restricted Boltzmann Machine (RBM) is made up of visible units, denoted as v_i , and hidden units, represented as h_μ . These units engage with one another through interactions characterized by the weights $w_{i\mu}$. Notably, there are no direct interactions among the visible units or among the hidden units themselves.

The goodness of a model is evaluated by computing the log-likelihood \mathcal{L} of the data:

$$\mathcal{L} = \frac{1}{M} \sum_{m=1}^M \ell_{\theta}(v^{(m)}) \quad (3)$$

$$\ell_{\theta}(v^{(m)}) = \log \sum_h e^{-E(v,h)} - \log Z \quad (4)$$

where M is the number of data points, Z is the partition function, $E(v, h)$ in the exponential argument $e^{-E(v, h)} = G(h) \prod_i e^{H_i(h)v_i}$ is the energy function, with $G(h) = \prod_\mu e^{b_\mu h_\mu}$ and $H_i(h) = a_i + \sum_\mu w_{i\mu} h_\mu$, and, in Eq. 4, the index h of the summation represents each possible state of the hidden layer.

3. METHODS

3.1. Theoretical implementation

Regarding the computation of the partition function Z (in Eq. 4), its implementation depends on the set of values which each unit can be set on. Since the direct computation of the Z function is unfeasible, due to the summation of all possible state combinations of the system, a possible approach is to split the summation separately over states v and h .

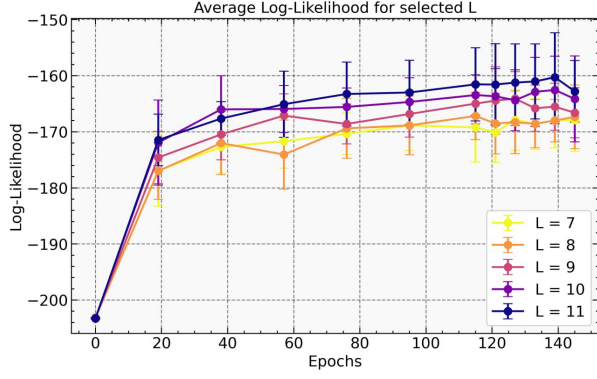


FIG. 2. The evolution of the log-likelihood over training epochs for different values of L is shown. The curves represent the average log-likelihood across 10 stochastically distinct runs.

For the Bernulli case the units can assume the values $\{0, 1\}$ and the partition function form can be manipulated to obtain the following expression:

$$Z = \sum_v \sum_h e^{-E(v,h)} = q^D \sum_h G(h) \prod_{i=1}^D \frac{1 + e^{H_i(h)}}{q} \quad (5)$$

where q is the average of $1 + e^{H_i(h)}$ $\forall i, h$ that we introduce to avoid overflow. Concerning the training process of RBM, v_i are set using real data. Weights $w_{i\mu}$ are initialized sampling values from a Gaussian distribution of mean 0 and standard deviation of $2/\sqrt{L + D}$. Biases b_μ of hidden units are initialized to 0. Thus h_μ are evaluated by Eq. 1 before passed to Eq. 2 in order to compute back v_i . In such evaluation it is convenient to set $a_i = \log[t_i/(1 - t_i)]$; fixed the i -th visible unit, t_i is defined as the number of times such unit is 1 over the whole training array set, normalized over the size of the set. Referring to Eq. 2, the choice of shifting the sigmoid $\sigma(x)$ by the function $\log[t_i/(1 - t_i)]$ of the average t_i ensures that the units can activate properly, even given initial weights $w_{i\mu}$ close to 0.

The initialization of biases a_i , firstly proposed by Hinton [2], makes hidden units able to activate and differentiate better their activation based on data, avoiding the risk of getting stuck on values near to 0.

3.2. Computation

We define two Python3 classes, one for the computation of the log-likelihood and the other for defining the RBM structure and managing all the input and output of the training process. Referring to Eq. 4, in order to address the problem of overflow of $e^{H_i(h)v_i}$, we approximate the summation of powers over h as a sum of the biggest ones. The purpose of this choice is to




FIG. 3. Bar chart displaying the weighted mean and standard deviation of the log-likelihood over the last epochs for different values of L .

extract the biggest term out of the summation thus the applied logarithm avoids exponentiation.

Regarding the expression of the partition function shown in Eq. 5, it changes form in Spin case, since it is characterized by units taking values from $\{-1, 1\}$; manipulating the second member of the equation one gets:

$$Z = \sum_v \sum_h e^{-E(v,h)} = \sum_h G(h) \prod_{i=1}^D 2 \cosh(H_i(h)) \quad (6)$$

For Eq. 6 the $2 \cosh(H_i(h))$ term is approximated to just $e^{H_i(h)}$ by adding a common positive value to all the exponents, which shifts them away from negative domain, and compensating that out of the summation. Thus the same approximation procedure described for Eq. 4 is applied to Eq. 6.

To upload the values of the weights and biases during the training phase, we adopted two different gradient descent types: SGD and RMSprop, whose update equations are found in the Appendix (Eq. 8 and Eq. 9). We also make use of the LASSO regularization (?) controlled by the constant γ (see Eq. 7).

We define a function to visualize the evolution of the log-likelihood of a specific model throughout the training phase and another function to compute and display the mean log-likelihood and standard deviation over the last few epochs.

For our analysis, we select the set of digits $\{0, 1, 2\}$ from the MNIST database. We then evaluate the performance of the RBM by utilizing the visualization functions to compare different configurations of hyperparameters. In order to account for stochasticity, for each chosen set of hyperparameters we first perform ten independent training runs varying the NumPy random seed and then compute the mean log-likelihood accross runs obtaining an average trend. To derive a final log-likelihood value for each hyperparameter configuration, we calculate a weighted average of the values of the log-likelihood over

ZZZZZZZZZZZZZZZZZZZZZ ZZZZZ ZZZZZZZZZZZZ ZZZZZZZZ Z ZZZZZZZZZZZZZ
 ZZZZ ZZZZZZZZZZZ ZZZZZZZZZZZZZZZZZZZ ZZZZZ ZZZZZZZZZZZZ ZZZZZZZZ
 Z ZZZZZZZZZZZZ ZZZZ ZZZZZZZZZZZ ZZZZZZZZZZZZZZZZZZZZZ ZZZZZ
 ZZZZZZZZZZZZ ZZZZZZZZ Z ZZZZZZZZZZZZZ ZZZZ ZZZZZZZZZZZ
 ZZZZZZZZZZZZZZZZZZZZZ ZZZZZ ZZZZZZZZZZZZ ZZZZZZZZ Z ZZZZZZZZZZZZZ
 ZZZZ ZZZZZZZZZZZ ZZZZZZZZZZZZZZZZZZZZZ ZZZZZ ZZZZZZZZZZZZ ZZZZZZZZ Z

zzzzzzzzzzzz zzzz zzzzzzzzzz.

5. CONCLUSIONS

Discuss the key aspects that we can take home from this work.

Check if your text is light, swift, and correct in exposing its passages.

APPENDIX

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \gamma \cdot \eta_t \cdot \text{sign}(\boldsymbol{\theta}) \quad (7)$$

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \eta_t \nabla_{\boldsymbol{\theta}} E(\boldsymbol{\theta}) \quad (8)$$

$$\begin{cases} \mathbf{g}_t = \nabla_{\boldsymbol{\theta}} E(\boldsymbol{\theta}) \\ \mathbf{s}_t = \beta \mathbf{s}_{t-1} + (1 - \beta) \mathbf{g}_t^2 \\ \boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \eta_t \frac{\mathbf{g}_t}{\sqrt{\mathbf{s}_t + \epsilon}} \end{cases} \quad (9)$$

-
- [1] Pankaj Mehta, Marin Bukov, Ching-Hao Wang, Alexandre G.R. Day, Clint Richardson, Charles K. Fisher, David J. Schwab, *A high-bias, low-variance introduction to Machine Learning for physicists*, Physics Reports **810**, 1–124 (2019).
 - [2] Geoffrey E. Hinton, *A Practical Guide to Training Restricted Boltzmann Machines*, Neural Networks: Tricks of the Trade: Second Edition, 599–619 (2012).

Assignment score grid

Structure: the exposition follow a logic order	8
Clarity: the text is brief enough, avoids complicated sentences and specifies all concepts and links	8
Depth: the text is not a shallow repetition of notions, there emerges a good understanding	8
Rigor: the analysis of the results is precise, quantitatively, and convincing	8
Innovation: new methods/ideas are introduced; conclusions beyond what introduced in the class	4