

GROUP2505 – Log-likelihood analysis for Restricted Boltzmann Machines

Elias Maria Bonasera, Alberto Casellato, Nicola Garbin, and Francesco Pazzocco
(Dated: March 26, 2025)

Here, in the “abstract” (more or less of 8 lines), there should be a short summary of the work and of its main findings. In a paper, the abstract is important also for attracting potential readers, hence it is convenient to start it with some catchy statement. — The rest of this text has the double purpose of (a) providing a template for the assignment latex, and (b) introducing how to structure the backbone of the text and explaining some details of this assignment. The “zzz” fill the space to show a typical (but not rigid) extension of the parts.

ZZZZZZZZZZZZZZ ZZZZ ZZZZZZZZZZ ZZZZZZZ
Z ZZZZZZZZZZZZ ZZZZ ZZZZZZZZZZ ZZZZZZZZZZZZZ ZZZZZ ZZZZZZZZZZZ ZZZZZZZ Z ZZZZZZZZZZZ ZZZZ ZZZZZZZZZZ
ZZZZZZZZZZZZZZZZZZ ZZZZZ ZZZZZZZZZZZ ZZZZZZZZ Z ZZZZZZZZZZZ ZZZZ ZZZZZZZZZZZ
ZZZZZZZZZZZZZZZZZZ ZZZZZ ZZZZZZZZZZZ ZZZZZZZZ Z ZZZZZZZZZZZ ZZZZ ZZZ

1. INTRODUCTION

The main goal of this assignment is to explore how the performance of an RBM changes for different choices of the hyperparameters of the model, using the MNIST digits as the database; in particular using the log-likelihood evaluation we explore the trend of the model during the learning process as the number of epochs increases.

[illegible]

2. RESTRICTED BOLTZMANN MACHINES

Restricted Boltzmann Machines (RBM) are a powerful kind of generative models designed to accomplish training processes relatively fast. In RBMs, a set of D binary visible units i of state v_i is symmetrically connected to a set of L binary hidden units μ of state h_μ ; the continuous weight $w_{i\mu}$ quantifies the strength between unit i and unit μ (see Figure 1). RBMs use an energy function to define the probability distribution over the input data. In the training process the energy of the configuration is minimized by adjusting the parameters θ . The most common training algorithm is contrastive divergence which allows to approximate the gradient of the likelihood to update the parameters. During this process, a cyclic Gibbs sampling is performed setting the visible units given the hidden ones and vice versa, according to the following probabilities:

$$p(h_\mu = 1 \mid \mathbf{v}) = \sigma(b_\mu + \sum_i v_i w_{i\mu}) \quad (.1)$$

$$p(v_i = 1 \mid \mathbf{h}) = \sigma(a_i + \sum_{\mu} h_{\mu} w_{i\mu}) \quad (.2)$$




FIG. 1. A Restricted Boltzmann Machine (RBM) is made up of visible units, denoted as v_i , and hidden units, represented as h_μ . These units engage with one another through interactions characterized by the weights $W_{i\mu}$. Notably, there are no direct interactions among the visible units or among the hidden units themselves.

where $\sigma(x) = 1/(1 + e^{-x})$ is the logistic sigmoid function, a_i is the bias of the i -th visible unit and b_μ is the bias of the μ -th hidden unit; they act shifting the sigmoid function $\sigma(x)$. The absence of links among units of the same type simplifies the training process. Moreover, the number of iterations of Eq. .1 and Eq. .2 can be setted to 1 if real data is used to fix \mathbf{v} in the first place.

The goodness of the models is evaluated by computing the log-likelihood \mathcal{L} of the data:

$$\mathcal{L} = \frac{1}{M} \sum_{m=1}^M \mathfrak{l}_{\theta}(v^{(m)}) \quad (.3)$$

$$\mathfrak{l}_\theta(v^{(m)}) = \log \sum_h e^{-E(v,h)} - \log Z \quad (.4)$$

where M is the number of data points, Z is the partition function, $e^{-E(v,h)} = G(h) \prod_i e^{H_i(h)v_i}$ is the energy function, with $G(h) = \prod_\mu e^{b_\mu h_\mu}$ and $H_i(h)$ defined as $H_i(h) = a_i + \sum_\mu w_{i\mu} h_\mu$, and, in Eq. 4, the index h of the summation represents each possible state of the hidden layer.

2. METHODS

2.1. Theoretical implementation

Regarding the computation of the partition function Z (in Eq. .4), its implementation depends on the set of values which each unit can be set on. Since the direct computation of the Z function is unfeasible, due to the summation of on all possible state combinations of the system, a possible approach is to split the summation separately over v and h .

For the Bernulli case the units can assume the values $\{0, 1\}$ and the partition function form can be manipulated to obtain the following expression:

$$Z = \sum_v \sum_h e^{-E(v,h)} = q^D \sum_h G(h) \prod_{i=1}^D \frac{1 + e^{H_i(h)}}{q} \quad (.5)$$

where q is the average of $1 + e^{H_i(h)}$ $\forall i, h$ that we introduce to avoid overflow. Concerning the training process of RBM, \mathbf{v} is set using real data. Weights $w_{i\mu}$ are initialized sampling values from a Gaussian distribution of mean 0 and standard deviation of $2/\sqrt{L + D}$. Biases b_μ of hidden units are initialized to 0. Thus \mathbf{h} is evaluated by Eq. .1 before passed to Eq. .2 in order to compute back \mathbf{v} . In such evaluation it is convenient to set $a_i = \log[p_i/(1 - p_i)]$; fixed the i -th visible unit, p_i is defined as the number of times such unit is on over the whole training array set, normalized over the size of the set. Referring to Eq. .2, the choice of shifting the sigmoid $\sigma(x)$ by the function $\log[p_i/(1 - p_i)]$ of the average p_i ensures that the units can activate properly, even given initial weights $w_{i\mu}$ close to 0.

The initialization of biases a_i , firstly proposed by Hinton, makes hidden units able to activate and differentiate better their activation based on data, avoiding the risk of getting stuck on values near to 0.

2.2. Computation

We define two Python3 classes, one for the computation of the log-likelihood and the other for defining the RBM structure and managing all the input and output of the training process. Referring to Eq. .4, in order to address the problem of overflow of $e^{H_i(h)v_i}$, we approximate the summation of powers over h as a sum of the biggest ones. The purpose of this choice is to extract the biggest term out of the summation.

Regarding the expression of the partition function shown in Eq. .5, it changes form for Spin case, characterized by units taking values from $\{-1, 1\}$. Manipulating the second member of the equation one gets:

$$Z = \sum_v \sum_h e^{-E(v,h)} = \sum_h G(h) \prod_{i=1}^D 2 \cosh(H_i(h)) \quad (.6)$$

For Eq. .6 the $2 \cosh(H_i(h))$ term is approximated to just $e^{H_i(h)}$ by shifting all exponents away from negative values by adding a positive value and properly compensating that out of the summation. Thus the same approximation procedure described for Eq. .5 is applied to Eq. .6.

To upload the values of the weights and biases during the training phase, we adopted two different gradient descent types: SGD and RMSprop, whose update equations are found in the Appendix (Eq. .8 and Eq. .9). We also make use of the LASSO regularization controlled by the constant γ (see Eq. .7).

We define a function to visualize the evolution of the log-likelihood of a specific model throughout the training phase and another function to compute and display the mean log-likelihood and standard deviation over the last few epochs.

For our analysis, we select the set of digits $\{0, 1, 2\}$ from the MNIST database. We then evaluate the performance of the RBM by utilizing the visualization functions to compare different configurations of hyperparameters using pre-trained models. To better highlight the stochasticity of the process, we first run the exact same model five times. We then specifically examine the influence of variations in the number of Contrastive Divergence steps, in the number of hidden units, in parameters such as the learning rate η , ϵ and γ . Additionally, we analyze the impact of the optimizer choice for models with one-hot encoding enabled (POTTS) or disabled and either Bernoulli or Spin variables.

3. RESULTS

zzzzzzzzzzzzzzzzzz zzzzz zzzzzzzzzzz zzzzzzz z zzzzzzzzzzz
zzzz zzzzzzzzzzz zzzzzzzzzzzzzzzzz zzzzz zzzzzzzzzzz zzzzzzz
z zzzzzzzzzzzzz zzzzz zzzzzzzzzzz zzzzzzzzzzzzzzzzz zzzzz
zzzzzzzzzzzz zzzzzzz z zzzzzzzzzzzzz zzzzz zzzzzzzzzzz
zzzzzzzzzzzzzzzzzz zzzzz zzzzzzzzzzz zzzzzzz z zzzzzzzzzzz
zzzz zzzzzzzzzzz zzzzzzzzzzzzzzzzz zzzzz zzzzzzzzzzz zzzzzzz z
zzzzzzzzzzzz zzzzz zzzzzzzzzzz.

zzzzzzzzzzzzzzzzzz zzzzz zzzzzzzzzzz zzzzzzz z zzzzzzzzzzz
zzzz zzzzzzzzzzz zzzzzzzzzzzzzzzzz zzzzz zzzzzzzzzzz zzzzzzz
z zzzzzzzzzzzzz zzzzz zzzzzzzzzzz zzzzzzzzzzzzzzzzz zzzzz
zzzzzzzzzzzz zzzzzzz z zzzzzzzzzzzzz zzzzz zzzzzzzzzzz
zzzzzzzzzzzzzzzzzz zzzzz zzzzzzzzzzz zzzzzzz z zzzzzzzzzzz
zzzz zzzzzzzzzzz zzzzzzzzzzzzzzzzz zzzzz zzzzzzzzzzz zzzzzzz
z zzzzzzzzzzzzz zzzzz zzzzzzzzzzz. zzzzzzzzzzzzzzzzz zzzzz
zzzzzzzzzzzz zzzzzzz z zzzzzzzzzzzzz zzzzz zzzzzzzzzzz
zzzzzzzzzzzzzzzzzz zzzzz zzzzzzzzzzz zzzzzzz z zzzzzzzzzzz
zzzz zzzzzzzzzzz zzzzzzzzzzzzzzzzz zzzzz zzzzzzzzzzz zzzzzzz
z zzzzzzzzzzzzz zzzzz zzzzzzzzzzz zzzzzzzzzzzzzzzzz zzzzz

ZZZZZZZZZZZZZZZZ ZZZZZ ZZZZZZZZZZ ZZZZZZZ Z ZZZZZZZZZZZZZ
ZZZZ ZZZZZZZZZZ ZZZZZZZZZZZZZZZ ZZZZZ ZZZZZZZZZZZ ZZZZZZZ
Z ZZZZZZZZZZZ ZZZZ ZZZZZZZZZZ ZZZZZZZZZZZZZZZ ZZZZZZZ
ZZZZZZZZZZ ZZZZZZZ Z ZZZZZZZZZZZ ZZZZ ZZZZZZZZZ
ZZZZZZZZZZZZZZZZZZ ZZZZZ ZZZZZZZZZZZ ZZZZZZZ Z ZZZZZZZZZZZ
ZZZZ ZZZZZZZZZZ ZZZZZZZZZZZZZZZ ZZZZZ ZZZZZZZZZZZ ZZZZZZZ
Z ZZZZZZZZZZZ ZZZZ ZZZZZZZZZ. ZZZZZZZZZZZZZZZ ZZZZZ
ZZZZZZZZZZ ZZZZZZZ Z ZZZZZZZZZZZ ZZZZ ZZZZZZZZZ

ZZZZZZZZZZZZZZZZZZZZ ZZZZZZ ZZZZZZZZZZZZ ZZZZZZZZ Z ZZZZZZZZZZZZ
ZZZZZ ZZZZZZZZZZZZ ZZZZZZZZZZZZZZZZZZZZZ ZZZZZZ ZZZZZZZZZZZZ ZZZZZZZZ Z
ZZZZZZZZZZZZZZZZZZZZ ZZZZZZ ZZZZZZZZZZZZ.

APPENDIX

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \gamma \cdot \eta_t \cdot \text{sign}(\boldsymbol{\theta}) \quad (.7)$$

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \eta_t \nabla_{\boldsymbol{\theta}} E(\boldsymbol{\theta}) \quad (.8)$$

$$\begin{cases} \mathbf{g}_t = \nabla_{\boldsymbol{\theta}} E(\boldsymbol{\theta}) \\ \mathbf{s}_t = \beta \mathbf{s}_{t-1} + (1 - \beta) \mathbf{g}_t^2 \\ \boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \eta_t \frac{\mathbf{g}_t}{\sqrt{\mathbf{s}_t + \epsilon}} \end{cases} \quad (.9)$$

-
- [1] B. Franklin, J. Here There **10**, 20–40 (1800).
 [2] A. Einstein, Int. J. There Here **20**, 125–133 (1910).

Assignment score grid

Structure: the exposition follow a logic order	8
Clarity: the text is brief enough, avoids complicated sentences and specifies all concepts and links	8
Depth: the text is not a shallow repetition of notions, there emerges a good understanding	8
Rigor: the analysis of the results is precise, quantitatively, and convincing	8
Innovation: new methods/ideas are introduced; conclusions beyond what introduced in the class	4